

```

In [1]: import tensorflow as tf
        from tensorflow.keras.datasets import cifar10
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense
        from tensorflow.keras.utils import to_categorical
        import matplotlib.pyplot as plt
        import numpy as np
        from sklearn.metrics import classification_report, confusion_matrix

        # Load CIFAR-10 dataset
        (x_train, y_train), (x_test, y_test) = cifar10.load_data()

        # Normalize images (scaling pixel values between 0 and 1)
        x_train = x_train.astype('float32') / 255.0
        x_test = x_test.astype('float32') / 255.0

        # Convert labels to categorical (one-hot encoding)
        y_train = to_categorical(y_train, num_classes=10)
        y_test = to_categorical(y_test, num_classes=10)

        # Define class names
        class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

        # verify dataset shape
        print(f"Training data shape:{x_train.shape}")
        print(f"Test data shape:{x_test.shape}")

        # Build CNN model
        model = Sequential([
            Input(shape=(32,32,3)),
            Conv2D(32, (3, 3), activation='relu'),
            MaxPooling2D((2, 2)),
            Conv2D(64, (3, 3), activation='relu'),
            MaxPooling2D((2, 2)),
            Conv2D(64, (3, 3), activation='relu'),
            Flatten(),
            Dense(128, activation='relu'),
            Dropout(0.5),
            Dense(10, activation='softmax')
        ])
        model.summary()

        # Compile the model
        model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
        print(x_train.shape, y_train.shape)
        print(x_test.shape, y_test.shape)

        # Train the model
        history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_c

        # Evaluate the model
        test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)

```

```

print(f"Test Accuracy: {test_acc:.4f}")

# Predictions
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = np.argmax(y_test, axis=1)

# Classification Report
print("Classification Report:")
print(classification_report(y_true_classes, y_pred_classes, target_names=cla

# Confusion Matrix
conf_matrix = confusion_matrix(y_true_classes, y_pred_classes)
print("Confusion Matrix:\n", conf_matrix)

# Plot training loss and accuracy
plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Loss vs Epochs')

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy vs Epochs')

plt.show()

# Predict sample images
def plot_sample_predictions():
    fig, axes = plt.subplots(3, 5, figsize=(10, 6))
    axes = axes.ravel()
    for i in range(15):
        index = np.random.randint(0, len(x_test))
        axes[i].imshow(x_test[index])
        axes[i].set_title(f"Pred: {class_names[y_pred_classes[index]]}\nTrue
        axes[i].axis('off')
    plt.tight_layout()
    plt.show()

plot_sample_predictions()

```

Downloading data from [https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.g](https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz)
z

170498071/170498071 4s 0us/step

Training data shape:(50000, 32, 32, 3)

Test data shape:(10000, 32, 32, 3)

Model: "sequential"

Layer (type)	Output Shape
conv2d (Conv2D)	(None, 30, 30, 32)
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)
conv2d_1 (Conv2D)	(None, 13, 13, 64)
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)
conv2d_2 (Conv2D)	(None, 4, 4, 64)
flatten (Flatten)	(None, 1024)
dense (Dense)	(None, 128)
dropout (Dropout)	(None, 128)
dense_1 (Dense)	(None, 10)

Total params: 188,810 (737.54 KB)

Trainable params: 188,810 (737.54 KB)

Non-trainable params: 0 (0.00 B)

```

(50000, 32, 32, 3) (50000, 10)
(10000, 32, 32, 3) (10000, 10)
Epoch 1/10
782/782  64s 78ms/step - accuracy: 0.2863 - loss: 1.8959
- val_accuracy: 0.5223 - val_loss: 1.3139
Epoch 2/10
782/782  80s 76ms/step - accuracy: 0.5133 - loss: 1.3600
- val_accuracy: 0.6032 - val_loss: 1.1047
Epoch 3/10
782/782  59s 75ms/step - accuracy: 0.5843 - loss: 1.1781
- val_accuracy: 0.6162 - val_loss: 1.0747
Epoch 4/10
782/782  81s 75ms/step - accuracy: 0.6193 - loss: 1.0787
- val_accuracy: 0.6447 - val_loss: 0.9876
Epoch 5/10
782/782  59s 75ms/step - accuracy: 0.6518 - loss: 0.9963
- val_accuracy: 0.6831 - val_loss: 0.9005
Epoch 6/10
782/782  83s 77ms/step - accuracy: 0.6784 - loss: 0.9313
- val_accuracy: 0.6867 - val_loss: 0.8734
Epoch 7/10
782/782  86s 82ms/step - accuracy: 0.6922 - loss: 0.8810
- val_accuracy: 0.6979 - val_loss: 0.8699
Epoch 8/10
782/782  79s 78ms/step - accuracy: 0.7098 - loss: 0.8265
- val_accuracy: 0.7070 - val_loss: 0.8313
Epoch 9/10
782/782  80s 75ms/step - accuracy: 0.7218 - loss: 0.7913
- val_accuracy: 0.7141 - val_loss: 0.8289
Epoch 10/10
782/782  57s 73ms/step - accuracy: 0.7340 - loss: 0.7635
- val_accuracy: 0.7089 - val_loss: 0.8417
313/313 - 4s - 13ms/step - accuracy: 0.7089 - loss: 0.8417
Test Accuracy: 0.7089
313/313  3s 10ms/step

```

Classification Report:

	precision	recall	f1-score	support
airplane	0.78	0.68	0.73	1000
automobile	0.89	0.80	0.84	1000
bird	0.67	0.55	0.60	1000
cat	0.51	0.55	0.53	1000
deer	0.63	0.63	0.63	1000
dog	0.66	0.55	0.60	1000
frog	0.78	0.79	0.78	1000
horse	0.69	0.82	0.75	1000
ship	0.72	0.89	0.80	1000
truck	0.77	0.83	0.80	1000
accuracy			0.71	10000
macro avg	0.71	0.71	0.71	10000
weighted avg	0.71	0.71	0.71	10000

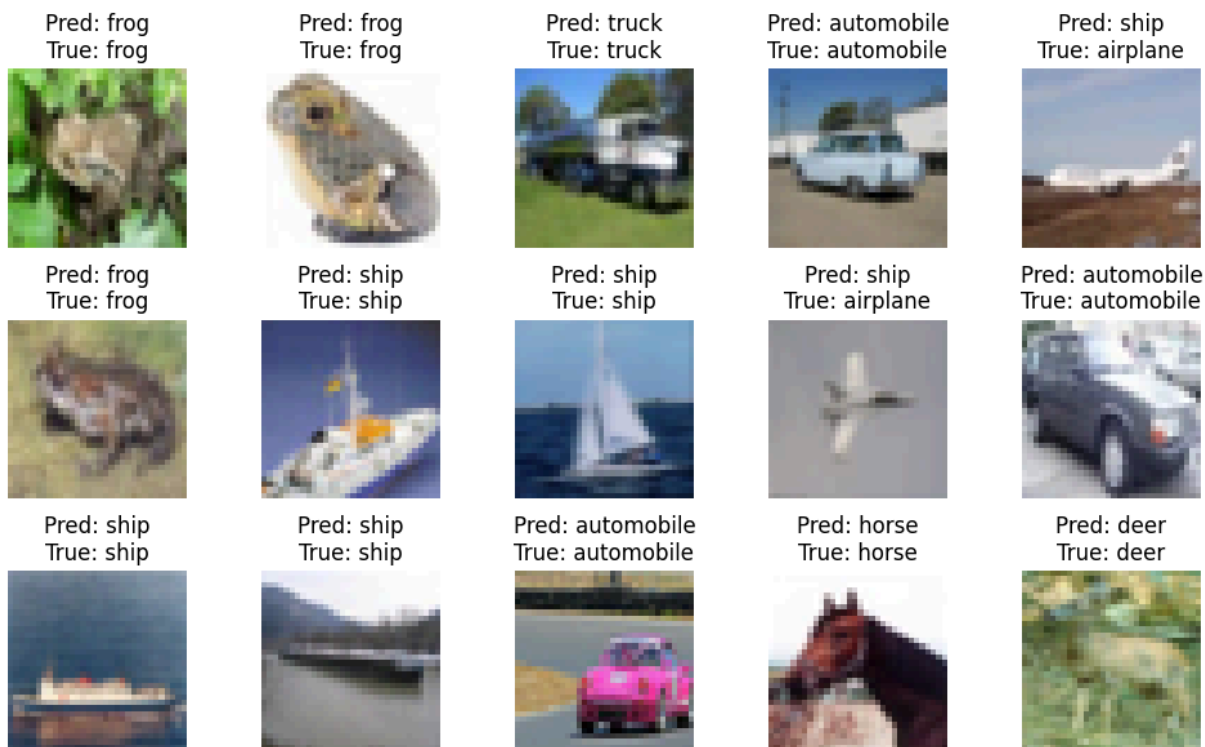


Confusion Matrix:

```

[[676 18 56 17 23 3 7 17 150 33]
 [ 21 707 4 18 3 1 6 7 35 108]

```

```
[ 57  2 551  76 116  53  64  46  19  16]
[ 18  7  56 545  81 145  58  41  28  21]
[ 15  2  51  71 630  13  60 131  24  3]
[  5  0  44 201  54 553  20  83  23 17]
[  6  4  22  72  47  16 794  13  15 11]
[ 12  2  25  37  43  39  4 816  7  15]
[ 31 16  8  18  2  4  3  4 894 20]
[ 21 45  6  15  1  5  7  18 49 833]]
```



This notebook was converted with convert.ploomber.io